

# Data Engineer Exercise

## 1. What considerations would you have for managing this data?

1. Points/questions that may affect solution design and architecture
  1. Main data use cases / number of users
  2. Types of data sources and targets
  3. Volume of the incoming and total data
  4. Velocity of data changes (stream, micro-batch, batch processing)
  5. Flexibility in toolset and primary/secondary cloud provider or internal data centers
  6. Security requirements
  7. Is there any specific legal regulation that has to be applied to the data (GDPR/CCPA, HIPAA, PCI etc)
2. The main challenges would be
  1. Educate end users to work with DWH engine of choice efficiently
  2. Keep costs under control simultaneously meeting all functional and non-functional SLA
  3. Build system of logging, monitoring and alerting that will guarantee work of application with acceptable error levels

### Some general thoughts on a high level architecture

1. All the raw files are saved perpetually in s3 data lake and can be pruned according to the corporate storage policies
  1. as a part of optimisation, source files could be converted to a store efficient formats like parquet or ORC with partitioning by day and then stored in s3
2. Apache Spark on Amazon EMR or Databricks Spark can be used as primary ELT tool, as an alternative or addition can be used AWS Data Pipelines, AWS Glue
  1. access data from s3, external SQL databases, Redshift and to
  2. cleanse, aggregate and send data to Redshift / s3
  3. specific pre-aggregated table depends on business needs, but from the example dataset we can build a fact table with number and duration for each kind of alert for the same 4 hour time window or 1 day window, with linkage to the detail table if user needs to go to the individual event level.
3. Amazon Redshift would be the main data warehouse engine that keeps internal data as well as providing access to federated sources in s3 via Amazon Redshift Spectrum, alternatively, consider Snowflake or even Firebolt on AWS new high performance and cost efficient Cloud Data Warehouse (firebolt.io)
4. Access configuration should follow the principle of least privilege.

1. majority of business users get read-only access only to the datasets that are required for them to complete tasks
2. group of power-users have read-write permissions to necessary datasets
3. and on the top of the pyramid are super users with full access and administrative permissions

## 2. Draft description for the device\_alerts view and each of its columns

### Assumptions

- This table represents stream of alerts each of them is signaling that device is online or offline
- Looking on the RSU\_ID pattern I can assume that potentially there could be billions of devices
- Type and geo location of each device could be coded in RSU\_ID
- To simplify analysis every event is saved in UTC time
- Alert event is active beginning from the alert time up to the receiving next alert.

### Table description

Column Name	Data Type	Description
timestamp_bigint	Long	Date/Time of alert in form of unix epoch with milliseconds
timestamp_utc	Timestamp(0)	Date/time of alert in UTC
rsu_id	String	Unique ID of rsu object
alert	String	Alert type {RSU_OFFLINE, POE_OFFLINE, SNMP_FAILURE, NONE}

### Alert type

- RSU\_OFFLINE - Road side unit is offline
- POE\_OFFLINE - Power-over-Ethernet failure
- SNMP\_FAILURE - Simple Network Management Protocol failure
- NONE - Road side unit works normally

## 3. What are the 4 hours in the time range with the most alerts triggered

### Simple 4 hour windowing

**4 hours between 2020/07/02 4pm and 8 pm had the most number of alerts equal to 56**

- Dataset already contains alert event in epoch format that is why we can easily split it on 4 hour buckets by deviding timestamp\_bigint on number of milliseconds in this 4 hours and discarding fractional part and than moving back to epoch time but without milliseconds as we do not need such precision

- Also we need to filter out alert type NONE as according to the assumptions it means that RSU is alive and not actually a health alert but rather signals that previous alert has completed
- After that we just need to count events in each 'window' and get the time with the largest result

```
+-----+-----+
|alert_hour      |cnt|rnk_val|
+-----+-----+
|2020-07-02 20:00:00|56 |1      |
+-----+-----+
```

```
select
  *
from
  (
    select
      alert_hour
    , cnt
    , rank(cnt) over (partition by rsu_id order by cnt desc) as rnk_val
    from
      (
        select
          from_unixtime(Decimal(int(timestamp_bigint/1000/60/60/4)) * 60 * 60 * 4,"yyyy-MM-dd HH
        , rsu_id
        , count(*) as cnt
        from device_alerts t1
        where alert != 'NONE'
        group by 1, 2
      )
  )
where rnk_val = 1
```

## Rolling 4 hour windowing

The other approach would be to count event on the rolling basis

- For each event we look 4 hours behind using self join
- As with previous approach all NONE alerts have to be filtered out
- Lastly counting number of alerts within 4 hour and take the largest results

**Using this approach we are getting 6 windows with 58 alerts each**

- 2 windows at 2020-07-02
- 1 window at 2020-07-17
- 3 windows at 2020-07-20

first_alert_ts	last_alert_ts	rsu_id	cnt_alerts	rnk_val
2020-07-02 19:46:12	2020-07-02 23:46:09	f943c318-26c8-45bz-zdzb-d43bz0e56690	58	1
2020-07-02 19:53:41	2020-07-02 23:51:13	f943c318-26c8-45bz-zdzb-d43bz0e56690	58	1
2020-07-17 12:19:17	2020-07-17 16:18:57	f943c318-26c8-45bz-zdzb-d43bz0e56690	58	1
2020-07-20 11:40:29	2020-07-20 15:40:09	f943c318-26c8-45bz-zdzb-d43bz0e56690	58	1
2020-07-20 11:44:20	2020-07-20 15:42:40	f943c318-26c8-45bz-zdzb-d43bz0e56690	58	1
2020-07-20 11:51:11	2020-07-20 15:50:12	f943c318-26c8-45bz-zdzb-d43bz0e56690	58	1

```

WITH vlt_alerts as
(
    SELECT
        timestamp_bigint
    , timestamp_utc
    , rsu_id
    FROM device_alerts
    WHERE alert != 'NONE'
)
SELECT
    *
FROM
    (
        SELECT
            *
        , RANK(cnt_alerts) OVER (PARTITION BY rsu_id ORDER BY cnt_alerts DESC) as rnk_val
    FROM
        (
            SELECT
                min(t1.timestamp_utc) as first_alert_ts
            , t.timestamp_utc as last_alert_ts
            , t.rsu_id
            , count(*) as cnt_alerts
            FROM vlt_alerts t, vlt_alerts t1
            WHERE t.rsu_id = t1.rsu_id
            AND t1.timestamp_bigint between t.timestamp_bigint - 4*60*60*1000 and t.timestamp_bigint
            group by 2, 3
        )
    )
WHERE rnk_val = 1
ORDER BY 1

```

# What trends do you notice between alert types?

1. After SNMP\_FAILURE always follow either POE\_OFFLINE or RSU\_OFFLINE
2. Overall alerts are not evenly distributed but are sporadically picking with long periods of healthy behaviour
3. SNMP\_FAILURE are usually the longest alert and overall took 44% of time having only 6% occurences
4. SNMP\_FAILURE and POE\_OFFLINE were happening toghether during the first half of the observing period, but in the second part SNMP\_FAILURE went to nothing

5. The most rare alert type is RSU\_OFFLINE with only 6 occurrences, that means that devices itself are more reliable than network

## What tools(s) did you choose to use and why?

1. Google Bigquery (in personal project) to run a few quick queries and get the basic understanding of the dataset
2. Databricks community Spark cluster running pySpark in Jupyter notebook for the main investigation as it is quite fast; interactive; provides convenient interfaces both in Dataframe and SQL

## 4. When did the alert with the longest duration occur?

alert	alert_begin_utc	alert_end_uts	alert_elapsed_time	rnk_val
SNMP_FAILURE	2020-07-03 01:56:14	2020-07-04 22:30:08	44 hours 33 minutes 54 seconds	1

```
select
  *
from
  (
    select
      alert
      , timestamp_utc as alert_begin_utc
      , end_timestamp_utc as alert_end_uts
      , end_timestamp_utc - timestamp_utc as alert_elapsed_time
      , rank(time_event) over (partition by rsu_id order by time_event desc) as rnk_val
    from
      (
        select
          t.alert
          , t.rsu_id
          , timestamp_utc
          , lead(timestamp_utc) over (partition by rsu_id order by timestamp_bigint) end_timestamp
          , lead(timestamp_bigint) over (partition by rsu_id order by timestamp_bigint) - timestamp
        from device_alerts t
      )
    where alert != 'NONE'
  )
where rnk_val = 1
```

## What tools(s) did you choose to use and why?

1. Spark cluster running SQL query on top of pySpark as I already had the table loaded there.

```
from pyspark.sql import SparkSession
from pyspark.sql.functions import *
from pyspark.sql.types import *

spark = (SparkSession
        .builder
        .appName("panasonic-assignment")
        .getOrCreate())

rsu_schema = StructType([StructField('timestamp_bigint', LongType(), True),
                           StructField('timestamp_utc', TimestampType(), True),
                           StructField('rsu_id', StringType(), True),
                           StructField('alert', StringType(), True)])

file_path = 'dbfs:/FileStore/shared_uploads/vasilii.surov@gmail.com/device_alerts.csv'
rsu_df = spark.read.csv(file_path, header=True, schema=rsu_schema)
rsu_df.cache()
rsu_df.printSchema()

root
|-- timestamp_bigint: long (nullable = true)
|-- timestamp_utc: timestamp (nullable = true)
|-- rsu_id: string (nullable = true)
|-- alert: string (nullable = true)

# How many rows do we have
rsu_df.count()

Out[3]: 3479

# Sample data
rsu_df.orderBy('timestamp_bigint').show(18, False)
```

timestamp_bigint	timestamp_utc	rsu_id	alert
1593583215452	2020-07-01 06:00:15	f943c318-26c8-45bz-zdzb-d43bz0e56690	POE_OFFLINE
1593583335091	2020-07-01 06:02:15	f943c318-26c8-45bz-zdzb-d43bz0e56690	NONE
1593583867309	2020-07-01 06:11:07	f943c318-26c8-45bz-zdzb-d43bz0e56690	SNMP_FAILURE
1593607123410	2020-07-01 12:38:43	f943c318-26c8-45bz-zdzb-d43bz0e56690	POE_OFFLINE
1593607170952	2020-07-01 12:39:30	f943c318-26c8-45bz-zdzb-d43bz0e56690	NONE
1593607574669	2020-07-01 12:46:14	f943c318-26c8-45bz-zdzb-d43bz0e56690	SNMP_FAILURE
1593619442468	2020-07-01 16:04:02	f943c318-26c8-45bz-zdzb-d43bz0e56690	POE_OFFLINE
1593619535172	2020-07-01 16:05:35	f943c318-26c8-45bz-zdzb-d43bz0e56690	NONE
1593619821456	2020-07-01 16:10:21	f943c318-26c8-45bz-zdzb-d43bz0e56690	POE_OFFLINE
1593619948069	2020-07-01 16:12:28	f943c318-26c8-45bz-zdzb-d43bz0e56690	NONE
1593620018361	2020-07-01 16:13:38	f943c318-26c8-45bz-zdzb-d43bz0e56690	POE_OFFLINE
1593620105083	2020-07-01 16:15:05	f943c318-26c8-45bz-zdzb-d43bz0e56690	NONE
1593620171361	2020-07-01 16:16:11	f943c318-26c8-45bz-zdzb-d43bz0e56690	POE_OFFLINE
1593620277183	2020-07-01 16:17:57	f943c318-26c8-45bz-zdzb-d43bz0e56690	NONE
1593620351447	2020-07-01 16:19:11	f943c318-26c8-45bz-zdzb-d43bz0e56690	POE_OFFLINE
1593620428113	2020-07-01 16:20:28	f943c318-26c8-45bz-zdzb-d43bz0e56690	NONE
1593620501477	2020-07-01 16:21:41	f943c318-26c8-45bz-zdzb-d43bz0e56690	POE_OFFLINE

```
# How many alert types do we have?
rsu_df.select('alert').distinct().show()
```

alert
RSU_OFFLINE
POE_OFFLINE
NONE
SNMP_FAILURE

```
# Double check that file contains data only for one object
rsu_df.select('rsu_id').distinct().show(truncate=False)
```

rsu_id
f943c318-26c8-45bz-zdzb-d43bz0e56690

Data Analysys

```
# Creating temporaty view
rsu_df.createOrReplaceTempView('device_alerts')
```

```
# Alert durations
spark.sql("""
select
    t.*
    , lead(timestamp_utc) over (partition by rsu_id order by timestamp_bigint) as end_timestamp_utc
    , lead(timestamp_utc) over (partition by rsu_id order by timestamp_bigint) - timestamp_utc as
elapsed_time
from device_alerts t
""").show(10, False)
```

timestamp_bigint	timestamp_utc	rsu_id	alert	end_timest
amp_utc	elapsed_time			
1593583215452	2020-07-01 06:00:15	f943c318-26c8-45bz-zdzb-d43bz0e56690	POE_OFFLINE	2020-07-01 06:02:15
2 minutes				
1593583335091	2020-07-01 06:02:15	f943c318-26c8-45bz-zdzb-d43bz0e56690	NONE	2020-07-01 06:11:07
8 minutes 52 seconds				
1593583867309	2020-07-01 06:11:07	f943c318-26c8-45bz-zdzb-d43bz0e56690	SNMP_FAILURE	2020-07-01 12:38:43
6 hours 27 minutes 36 seconds				
1593607123410	2020-07-01 12:38:43	f943c318-26c8-45bz-zdzb-d43bz0e56690	POE_OFFLINE	2020-07-01

12:39:30 47 seconds			
1593607170952	2020-07-01 12:39:30	f943c318-26c8-45bz-zdzb-d43bz0e56690	NONE   2020-07-01
12:46:14 6 minutes 44 seconds			
1593607574669	2020-07-01 12:46:14	f943c318-26c8-45bz-zdzb-d43bz0e56690	SNMP_FAILURE   2020-07-01
16:04:02 3 hours 17 minutes 48 seconds			
1593619442468	2020-07-01 16:04:02	f943c318-26c8-45bz-zdzb-d43bz0e56690	POE_OFFLINE   2020-07-01
16:05:35 1 minutes 33 seconds			
1593619535172	2020-07-01 16:05:35	f943c318-26c8-45bz-zdzb-d43bz0e56690	NONE   2020-07-01
16:10:21 4 minutes 46 seconds			
1593619821456	2020-07-01 16:10:21	f943c318-26c8-45bz-zdzb-d43bz0e56690	POE_OFFLINE   2020-07-01
16:12:28 2 minutes 7 seconds			
1593619948069	2020-07-01 16:12:28	f943c318-26c8-45bz-zdzb-d43bz0e56690	NONE   2020-07-01
16:13:38 1 minutes 10 seconds			

```
+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+
```

only showing top 10 rows

```
# dataframe with alert time bucketed to 4 hour window
df_hour = spark.sql("""
SELECT
    alert_hour
    , CASE WHEN alert = 'POE_OFFLINE' THEN alert_hour END as POE_OFFLINE
    , CASE WHEN alert = 'SNMP_FAILURE' THEN alert_hour END as SNMP_FAILURE
    , CASE WHEN alert = 'RSU_OFFLINE' THEN alert_hour END as RSU_OFFLINE
FROM
(
SELECT
    int(timestamp_bigint/1000/60/60/4) as alert_hour
    , alert
FROM device_alerts t1
WHERE alert != 'NONE'
)
""")
df_hour.cache()
df_hour.count()
```

Out[9]: 1855


```
!pip install pyspark_dist_explore
```

```
# Alerts distribution accros time line (all alerts together)
```

```
from pyspark_dist_explore import hist
import matplotlib.pyplot as plt

bins = df_hour.select('alert_hour').distinct().count()*3
fig, ax = plt.subplots()
fig.set_figheight(10)
fig.set_figwidth(20)
hist(ax, df_hour.select(col('alert_hour')), bins=bins)
labels = ['Total Alerts']
ax.legend(labels)
```






# Alerts distribution by type

```
fig, ax = plt.subplots()
fig.set_figheight(10)
fig.set_figwidth(20)
POE_OFFLINE = df_hour.select(col('POE_OFFLINE'))
SNMP_FAILURE = df_hour.select(col('SNMP_FAILURE'))
RSU_OFFLINE = df_hour.select(col('RSU_OFFLINE'))

bins = df_hour.select('alert_hour').distinct().count()*3

hist(ax, [POE_OFFLINE, SNMP_FAILURE, RSU_OFFLINE], bins=bins, overlapping=True)
labels = ['POE_OFFLINE', 'SNMP_FAILURE', 'RSU_OFFLINE']
ax.legend(labels)
ax.set_title('Alerts')
```



```
%sql
-- Duration by alert
SELECT
  alert
, SUM(alert_duration) as total_alert_duration
, COUNT(alert) as cnt_alert_duration
, int(avg(alert_duration)) as avg_alert_duration
FROM
  (
    SELECT
      t.alert
      , lead(timestamp_bigint) over (partition by rsu_id order by timestamp_bigint) -
timestamp_bigint as alert_duration
    FROM device_alerts t
  )
GROUP BY alert
```

	alert ▲	total_alert_duration ▲	cnt_alert_duration ▲	avg_alert_duration ▲
1	RSU_OFFLINE	34247	6	5707
2	POE_OFFLINE	274532401	1625	169047
3	NONE	1561377162	1624	961439
4	SNMP_FAILURE	1416684084	224	6324482

Showing all 4 rows.



# 4 hours in the time range

```
# Bucketing by 4 hour
spark.sql("""
select
  *
from
  (
    select
      alert_hour
    , cnt
    , rank(cnt) over (partition by rsu_id order by cnt desc) as rnk_val
    from
      (
        select
          from_unixtime(Decimal(int(timestamp_bigint/1000/60/60/4)) * 60 * 60 * 4,"yyyy-MM-dd
HH:mm:ss") as alert_hour
          , rsu_id
          , count(*) as cnt
        from device_alerts t1
        where alert != 'NONE'
        group by 1, 2
      )
    )
where rnk_val = 1""").show(truncate=False)
```

alert_hour	cnt	rnk_val
2020-07-02 20:00:00	56	1

```
# Rolling windows
spark.sql("""
WITH vlt_alerts as
(
    SELECT
        timestamp_bigint
    , timestamp_utc
    , rsu_id
    FROM device_alerts
    WHERE alert != 'NONE'
)
SELECT
    *
FROM
    (
        SELECT
            *
        , RANK(cnt_alerts) OVER (PARTITION BY rsu_id ORDER BY cnt_alerts DESC) as rnk_val
    FROM
        (
            SELECT
                min(t1.timestamp_utc) as first_alert_ts
            , t.timestamp_utc as last_alert_ts
            , t.rsu_id
            , count(*) as cnt_alerts
            FROM vlt_alerts t, vlt_alerts t1
            WHERE t.rsu_id = t1.rsu_id
            AND t1.timestamp_bigint between t.timestamp_bigint - 4*60*60*1000 and t.timestamp_bigint
            group by 2, 3
        )
    )
WHERE rnk_val = 1
ORDER BY 1
""").show(10, False)
```

first_alert_ts	last_alert_ts	rsu_id	cnt_alerts	rnk_val
2020-07-02 19:46:12	2020-07-02 23:46:09	f943c318-26c8-45bz-zdzb-d43bz0e56690	58	1
2020-07-02 19:53:41	2020-07-02 23:51:13	f943c318-26c8-45bz-zdzb-d43bz0e56690	58	1
2020-07-17 12:19:17	2020-07-17 16:18:57	f943c318-26c8-45bz-zdzb-d43bz0e56690	58	1
2020-07-20 11:40:29	2020-07-20 15:40:09	f943c318-26c8-45bz-zdzb-d43bz0e56690	58	1
2020-07-20 11:44:20	2020-07-20 15:42:40	f943c318-26c8-45bz-zdzb-d43bz0e56690	58	1
2020-07-20 11:51:11	2020-07-20 15:50:12	f943c318-26c8-45bz-zdzb-d43bz0e56690	58	1

## The longest duration

```
spark.sql("""
select
    *
from
    (
        select
            alert
            , timestamp_utc as alert_begin_utc
            , end_timestamp_utc as alert_end_uts
            , end_timestamp_utc - timestamp_utc as alert_elapsed_time
            , rank(time_event) over (partition by rsu_id order by time_event desc) as rnk_val
        from
            (
                select
                    t.alert
                    , t.rsu_id
                    , timestamp_utc
                    , lead(timestamp_utc) over (partition by rsu_id order by timestamp_bigint) end_timestamp_utc
                    , lead(timestamp_bigint) over (partition by rsu_id order by timestamp_bigint) -
timestamp_bigint as time_event
                from device_alerts t
            )
        where alert != 'NONE'
    )
where rnk_val = 1
""").show(truncate=False)
```

alert	alert_begin_utc	alert_end_uts	alert_elapsed_time	rnk_val
SNMP_FAILURE	2020-07-03 01:56:14	2020-07-04 22:30:08	44 hours 33 minutes 54 seconds	1

```
from pyspark.sql import Window
window = Window.orderBy("timestamp_bigint")
# rsu_df.orderBy('timestamp_bigint').show()
df2 = rsu_df.withColumn("next_alert", lead("alert").over(window)).withColumn("next_next_alert",
lead("alert", 2).over(window)).withColumn("next_next_next_alert", lead("alert", 3).over(window))
```

alert	next_alert	next_next_alert
NONE	POE_OFFLINE	null
NONE	POE_OFFLINE	NONE
NONE	RSU_OFFLINE	POE_OFFLINE
NONE	SNMP_FAILURE	POE_OFFLINE
NONE	SNMP_FAILURE	RSU_OFFLINE
POE_OFFLINE	null	null
POE_OFFLINE	NONE	POE_OFFLINE
POE_OFFLINE	NONE	RSU_OFFLINE
POE_OFFLINE	NONE	SNMP_FAILURE

	POE_OFFLINE	RSU_OFFLINE	NONE
	RSU_OFFLINE	NONE	POE_OFFLINE
	RSU_OFFLINE	POE_OFFLINE	NONE
	SNMP_FAILURE	POE_OFFLINE	NONE
	SNMP_FAILURE	POE_OFFLINE	RSU_OFFLINE
	SNMP_FAILURE	RSU_OFFLINE	POE_OFFLINE
+-----+-----+-----+			